

# Neural Networks from Math

Vu Hung Nguyen

2025-11-03

---

# Contents

<b>Preface</b>	iii
<b>1 Overview</b>	1
<b>2 Functions and Graphs</b>	3
2.1 From Tables to Rules . . . . .	3
2.2 Exercises . . . . .	4
<b>3 Data and Error</b>	5
3.1 Measuring Error . . . . .	5
3.2 Exercises . . . . .	6
<b>4 Vectors and Matrices</b>	7
4.1 Notation . . . . .	7
4.2 Exercises . . . . .	8
<b>5 Linear Models</b>	9
5.1 Form and Geometry . . . . .	9
5.2 Exercises . . . . .	9
<b>6 Nonlinearity and Activations</b>	11
6.1 Common Activations . . . . .	11
6.2 Exercises . . . . .	12
<b>7 Composing Layers</b>	13
7.1 Two-Layer Example . . . . .	13
7.2 Exercises . . . . .	14
<b>8 Loss and Optimisation</b>	15
8.1 Loss Choices . . . . .	15
8.2 Gradient Descent (Conceptually) . . . . .	16
<b>9 Training Loop Concepts</b>	18
9.1 The Loop . . . . .	18
9.2 Exercises . . . . .	19
<b>10 A Simple Neural Network Architecture</b>	20
10.1 Shapes and Flow . . . . .	20
10.2 Exercises . . . . .	21

<b>11 Limits and Ethics</b>	<b>22</b>
11.1 Limits . . . . .	22
11.2 Ethics . . . . .	22
<b>12 Glossary</b>	<b>24</b>
<b>A Worksheets: Checkpoints</b>	<b>25</b>
<b>B Worksheets: Exercises</b>	<b>26</b>
<b>References</b>	<b>27</b>
<b>About the Author</b>	<b>29</b>

---

# Preface

This brief book takes an example-first path from familiar school mathematics to the core ideas behind a simple neural network architecture. It is designed for motivated high-school students: short chapters, visual intuition, and concrete checkpoints.

# Chapter 1

## Overview

This book takes you from familiar mathematics to a working understanding of simple neural networks. It is written in clear British English, and designed for undergraduates and motivated high-school students.

### Learning Objectives

1. Understand the book's roadmap and how chapters connect.
2. Recognise the key ingredients of a neural network.
3. Learn how examples, metaphors, and visuals support intuition.

## Roadmap

We move from functions and graphs to data and error, then to vectors and matrices that compactly represent many inputs. Linear models give us a first predictive rule; nonlinear activations add the “kinks” that let us model more interesting patterns. Stacking layers composes these ideas. Finally, we discuss loss and optimisation, the training loop, a simple network, and practical limits and ethics.

### Example 1.1

Consider predicting house price from size. A function maps size (input) to price (output). A linear model draws a straight line; a neural network allows bends via activations and layers, improving fit when reality is not a straight line.

### Remark 1.1

Throughout, we prioritise intuition first, then formalism. Visuals accompany core definitions where helpful.

---

## How to Use This Book

Skim the learning objectives at the start of each chapter. Work through examples; attempt exercises before revealing the upside-down hints. Use the glossary to refresh key terms.

---

# Chapter 2

## Functions and Graphs

Functions describe how inputs map to outputs. Graphs help us see this mapping.

### Learning Objectives

1. Interpret a function as a rule from input to output.
2. Read tables and plots to understand trends.
3. Distinguish linear from nonlinear patterns visually.

### 2.1 From Tables to Rules

Suppose we record study time (hours) and score (percentage). A function  $f$  turns an input  $x$  into an output  $y = f(x)$ . A straight trend suggests a linear rule; bends suggest nonlinearity.

#### Definition 2.1

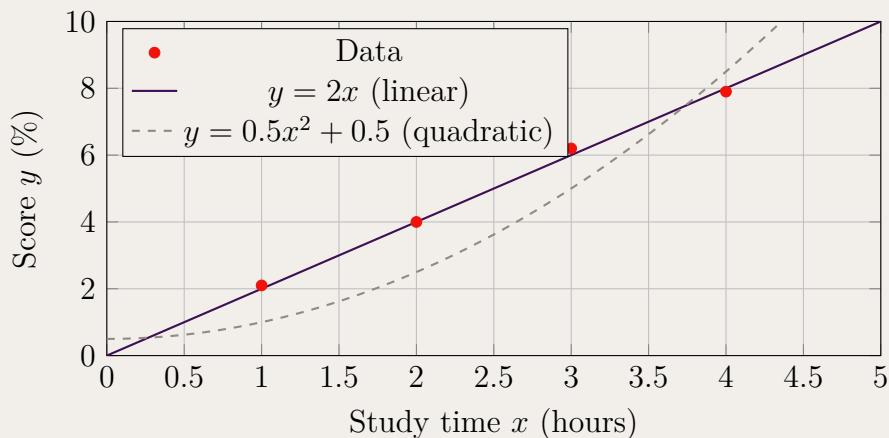
A *function* is a mapping assigning each input a single output. We often visualise  $(x, f(x))$  as points on a graph.

**Example 2.1**

Measured study time (hours) and score (%):

$x$	1	2	3	4
$y$	2.1	4.0	6.2	7.9

If doubling  $x$  roughly doubles  $y$ , a straight line is sensible. Below we plot the points, a simple linear rule  $y = 2x$ , and a curved (quadratic) rule  $y = 0.5x^2 + 0.5$  to contrast linear vs nonlinear behaviour.



The linear rule tracks the trend closely; the quadratic bends upward and will diverge from the roughly proportional pattern as  $x$  grows.

## 2.2 Exercises

Given points  $(1, 2)$ ,  $(2, 4.1)$ ,  $(3, 5.9)$ , would a straight line be a reasonable model? Explain briefly.

**Hint:** Turn the page upside down to read yes. The ratios are consistent with a near-linear trend; noise causes small deviations.

---

# Chapter 3

## Data and Error

Predictions meet reality through data. Error tells us how far off we are.

### Learning Objectives

1. Define prediction, target, and error for a dataset.
2. Compute simple average absolute and squared error.
3. Explain why averaging error stabilises noisy measurements.

### 3.1 Measuring Error

Given inputs  $x_i$  with observed outputs  $y_i$  and predictions  $\hat{y}_i$ , an error for item  $i$  is  $e_i = \hat{y}_i - y_i$ . Two popular summaries are the mean absolute error (MAE) and mean squared error (MSE):

$$\text{MAE} = \frac{1}{n} \sum_i |e_i|, \quad \text{MSE} = \frac{1}{n} \sum_i e_i^2.$$

**Example 3.1**

Using the study-score data from the previous chapter

$x$	1	2	3	4
$y$	2.1	4.0	6.2	7.9

consider the linear rule  $\hat{y} = 2x$ . The predictions and errors are:

$x$	1	2	3	4
$\hat{y} = 2x$	2.0	4.0	6.0	8.0
$e = \hat{y} - y$	0.1	0.0	-0.2	0.1
$ e $	0.1	0.0	0.2	0.1
$e^2$	0.01	0.00	0.04	0.01

Thus

$$\text{MAE} = \frac{0.1+0.0+0.2+0.1}{4} = 0.1, \quad \text{MSE} = \frac{0.01+0.00+0.04+0.01}{4} = 0.015.$$

The small MAE/MSE values indicate the line  $\hat{y} = 2x$  matches the trend closely for these points.

**Remark 3.1**

Squaring emphasises large mistakes; absolute value treats all deviations proportionally. Choice depends on the application.

## 3.2 Exercises

For true values  $y = (1, 2, 3)$  and predictions  $\hat{y} = (1.2, 1.9, 3.4)$ , compute MAE and MSE.

**Hint: Errors:**  $(0.2, -0.1, 0.4)$ .  $\text{MAE} = (0.2 + 0.1 + 0.4)/3 = 0.233 \dots$ ;  $\text{MSE} = (0.04 + 0.01 + 0.16)/3 = 0.07 \dots$

---

# Chapter 4

## Vectors and Matrices

Vectors collect features; matrices collect many vectors and define linear transformations.

### Learning Objectives

1. Represent data points as vectors of features.
2. Use matrix–vector multiplication to combine features linearly.
3. Interpret a matrix as a transformation of space.

### 4.1 Notation

Write a feature vector as  $\mathbf{x} = (x_1, \dots, x_d)^\top$ . A weight vector  $\mathbf{w}$  and bias  $b$  define a linear score  $s = \mathbf{w}^\top \mathbf{x} + b$ .

#### Definition 4.1

A *matrix*  $W \in \mathbb{R}^{m \times d}$  applied to a vector  $\mathbf{x} \in \mathbb{R}^d$  yields  $W\mathbf{x} \in \mathbb{R}^m$ , combining inputs into  $m$  outputs.

#### Example 4.1

With two features (height, width) and two outputs (sum, difference),  $W = \begin{bmatrix} 1 & 1 \\ 1 & -1 \end{bmatrix}$  maps a vector  $\mathbf{x} = (x_1, x_2)^\top$  to  $W\mathbf{x} = (x_1 + x_2, x_1 - x_2)^\top$ .  
For a concrete calculation, let  $\mathbf{x} = (3, 2)^\top$ . Then

$$W\mathbf{x} = \begin{bmatrix} 1 & 1 \\ 1 & -1 \end{bmatrix} \begin{bmatrix} 3 \\ 2 \end{bmatrix} = \begin{bmatrix} 1 \cdot 3 + 1 \cdot 2 \\ 1 \cdot 3 + (-1) \cdot 2 \end{bmatrix} = \begin{bmatrix} 5 \\ 1 \end{bmatrix}.$$

So the first output (sum) is 5 and the second output (difference) is 1.

## 4.2 Exercises

If  $W = \begin{bmatrix} 2 & 0 \\ 0 & 3 \end{bmatrix}$  and  $\mathbf{x} = (1, 2)^\top$ , compute  $W\mathbf{x}$ .

**Hint:**  $(2, 6)^\top$ .

---

# Chapter 5

## Linear Models

A linear model scores inputs by a weighted sum plus a bias; decisions follow from thresholds or regression.

### Learning Objectives

1. Write a linear model as  $\hat{y} = \mathbf{w}^\top \mathbf{x} + b$ .
2. Explain the geometric view: lines and hyperplanes.
3. Recognise when linear models are insufficient.

### 5.1 Form and Geometry

The set of points with constant score forms a hyperplane. When classes are linearly separable, a single hyperplane can divide them.

#### Example 5.1

Classify points above a line: take  $\hat{y} = 2x_1 - x_2 + 0.5$  and predict positive when  $\hat{y} > 0$ .

For  $(x_1, x_2) = (2, 1)$ ,

$$\hat{y} = 2 \cdot 2 - 1 + 0.5 = 3.5 > 0 \Rightarrow \text{positive.}$$

For  $(x_1, x_2) = (0, 1)$ ,

$$\hat{y} = 2 \cdot 0 - 1 + 0.5 = -0.5 < 0 \Rightarrow \text{negative.}$$

The decision boundary  $\hat{y} = 0$  is the straight line  $x_2 = 2x_1 + 0.5$ .

### 5.2 Exercises

Find a line that separates points A:  $(0, 2), (1, 3)$  from B:  $(2, 0), (3, 1)$  if possible.

**Hint:** Try  $x_2 = x_1 + 1.5$ . Points A lie above, B below.

# Chapter 6

## Nonlinearity and Activations

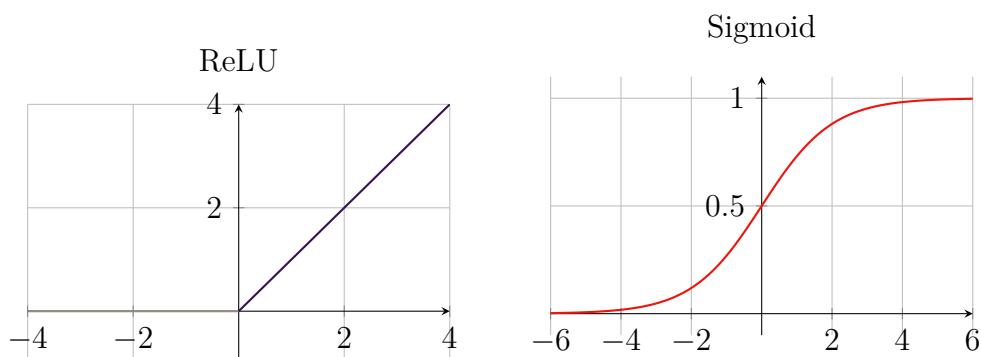
Nonlinear activations introduce bends that let models capture curved patterns.

### Learning Objectives

1. Describe common activations: ReLU, sigmoid, tanh.
2. Explain why stacking linear layers without activations stays linear.
3. Match activations to tasks (e.g., sigmoid for probabilities).

### 6.1 Common Activations

ReLU:  $\text{ReLU}(z) = \max(0, z)$  adds a kink at zero. Sigmoid  $\sigma(z) = 1/(1 + e^{-z})$  squashes to  $(0, 1)$ . Tanh squashes to  $(-1, 1)$ .



**Why this matters.** Activations create nonlinearity so networks can approximate curved relationships. ReLU is simple and keeps gradients flowing for positive inputs, aiding optimisation. Sigmoid maps real numbers to  $(0, 1)$ , ideal for probabilities; but it can saturate (flat tails), slowing learning when inputs are very large in magnitude.

**Bio example (sigmoid).** The logistic (sigmoid) curve models population growth with carrying capacity and neuron firing rates as a function of membrane potential: response is low for small input, steep near threshold, and saturates at high input.

**Remark 6.1**

Linear  $\circ$  linear is still linear. Nonlinearity between linear steps is essential.

## 6.2 Exercises

Sketch and compare  $\tanh(z)$  and leaky ReLU  $\text{LReLU}_\alpha(z) = \max(\alpha z, z)$  with  $\alpha = 0.1$ . Where are they most/least sensitive, and what ranges do their outputs take?

**Hint:** ReLU is flat for  $z < 0$  and slope 1 for  $z > 0$ . Sigmoid is steepest near 0, flat in tails.

---

# Chapter 7

## Composing Layers

Stack linear steps with activations to form complex functions.

### Learning Objectives

1. Define a layer as  $\mathbf{h} = g(W\mathbf{x} + \mathbf{b})$ .
2. Explain how two layers can form piecewise linear curves.
3. Understand hidden units as learned features.

### 7.1 Two-Layer Example

Let  $\mathbf{h} = \text{ReLU}(W_1\mathbf{x} + \mathbf{b}_1)$  and  $\hat{y} = W_2\mathbf{h} + b_2$ . Even with ReLU, the output is a flexible piecewise linear function of  $\mathbf{x}$ .

**Example 7.1**

In one dimension, two ReLUs can create a “bump” by combining kinks at different locations. Consider a tiny network with one input  $x$ , two hidden units, and one output:

$$\mathbf{h} = \text{ReLU}(W_1x + \mathbf{b}_1), \quad \hat{y} = W_2 \mathbf{h} + b_2,$$

where

$$W_1 = \begin{bmatrix} 1 \\ -1 \end{bmatrix}, \quad \mathbf{b}_1 = \begin{bmatrix} -1 \\ -3 \end{bmatrix}, \quad W_2 = \begin{bmatrix} 1 & -1 \end{bmatrix}, \quad b_2 = 0.$$

Thus the hidden units are

$$h_1 = \text{ReLU}(x - 1), \quad h_2 = \text{ReLU}(-x - 3) = \text{ReLU}(-(x + 3)).$$

The output becomes  $\hat{y} = h_1 - h_2$ . Evaluate at a few inputs:

$x$	$h_1 = \max(0, x - 1)$	$h_2 = \max(0, -x - 3)$	$\hat{y} = h_1 - h_2$
-5	0	2	-2
-3	0	0	0
0	0	0	0
2	1	0	1
4	3	0	3

Interpretation:

- For  $x \leq -3$ , only the second unit is active, so  $\hat{y} = -h_2$  decreases linearly.
- For  $-3 < x < 1$ , both units are off, so  $\hat{y} = 0$ .
- For  $x \geq 1$ , only the first unit is active, so  $\hat{y} = h_1$  increases linearly.

This piecewise linear shape forms a flat region with rising and falling flanks—one way to build a “bump” by composing ReLUs.

## 7.2 Exercises

How many linear regions can a sum of two shifted ReLUs create on the line?

**Hint:** Up to three regions separated by the two kink locations.

# Chapter 8

## Loss and Optimisation

Loss quantifies mismatch between predictions and targets; optimisation reduces loss.

### Learning Objectives

1. Define a loss function suitable for a task.
2. Describe gradient descent at a high level.
3. Relate learning rate to step size and stability.

### 8.1 Loss Choices

For regression, MSE is common; for binary classification, logistic loss pairs well with sigmoid outputs.

#### Example 8.1: Regression loss (MSE)

Suppose a model predicts house prices (in thousands) for three homes:  $\hat{y} = (210, 195, 250)$ , and true prices are  $y = (200, 205, 240)$ . Errors are  $e = \hat{y} - y = (10, -10, 10)$ . The mean squared error is

$$\text{MSE} = \frac{1}{3}(10^2 + (-10)^2 + 10^2) = \frac{1}{3}(100 + 100 + 100) = 100.$$

Units are squared (here, thousands<sup>2</sup>). A smaller value indicates closer predictions.

**Example 8.2: Logistic loss**

For a binary label  $y \in \{0, 1\}$  with sigmoid output  $p = \sigma(z)$  interpreted as  $\Pr(y = 1 | x)$ , the logistic loss for a single example is

$$\ell(y, p) = -\left(y \log p + (1 - y) \log(1 - p)\right).$$

Consider two cases with  $p = 0.9$ :

- If  $y = 1$ :  $\ell = -(1 \cdot \log 0.9 + 0 \cdot \log 0.1) \approx 0.105$  (small penalty).
- If  $y = 0$ :  $\ell = -(0 \cdot \log 0.9 + 1 \cdot \log 0.1) \approx 2.303$  (large penalty for confident wrong prediction).

This asymmetry encourages calibrated probabilities: confident and correct is rewarded; confident and wrong is penalised heavily.

## 8.2 Gradient Descent (Conceptually)

Imagine standing on a landscape where height is loss. At each step, move a little in the downhill direction; repeat until you are low enough.

**Example 8.3: One variable**

Let  $f(x) = (x - 3)^2$ . Its derivative is  $f'(x) = 2(x - 3)$ . Starting at  $x_0 = 0$  with learning rate  $\eta = 0.5$ , gradient descent updates are

$$x_{k+1} = x_k - \eta f'(x_k) = x_k - 0.5 \cdot 2(x_k - 3) = 3 - (x_k - 3).$$

Numerically:  $x_1 = 1.5$ ,  $x_2 = 2.25$ ,  $x_3 = 2.625\dots$  which approaches the minimiser  $x^* = 3$ .

**Example 8.4: Two variables**

Let  $f(x, y) = (x - 1)^2 + (y + 2)^2$ . The gradient is  $\nabla f = (2(x - 1), 2(y + 2))$ . With  $(x_0, y_0) = (0, 0)$  and  $\eta = 0.25$ :

$$(x_1, y_1) = (0, 0) - 0.25 (2(-1), 2(2)) = (0.5, -1).$$

Next step uses the new gradient:  $\nabla f(x_1, y_1) = (2(-0.5), 2(1)) = (-1, 2)$ , thus

$$(x_2, y_2) = (0.5, -1) - 0.25 (-1, 2) = (0.75, -1.5).$$

The iterates head toward the minimiser  $(1, -2)$ .

**Many inputs (neural nets).** For networks, parameters form a long vector  $\boldsymbol{\theta}$  (all weights and biases). The loss  $L(\boldsymbol{\theta})$  is averaged over a mini-batch. We compute the

gradient  $\nabla L(\boldsymbol{\theta})$  efficiently by backpropagation and update

$$\boldsymbol{\theta}_{k+1} = \boldsymbol{\theta}_k - \eta \nabla L(\boldsymbol{\theta}_k).$$

Different layers receive different components of the same gradient, moving all parameters a little in the direction that most reduces loss on the current batch.

If steps are too large, what failure can occur?

**Hint:** You can overshoot and bounce around (diverge) instead of settling.

---

# Chapter 9

## Training Loop Concepts

Predict, measure, adjust, repeat. Split data to measure progress fairly.

### Learning Objectives

1. Describe the predict–loss–update loop.
2. Explain train/validation/test splits.
3. Recognise overfitting and how validation helps.

### 9.1 The Loop

For each mini-batch: compute predictions, compute loss, update parameters a little. After each epoch, check validation loss.

**Example 9.1**

Consider a simple linear model  $\hat{y} = wx + b$  trained with mean squared error on one mini-batch of two points:  $(x, y) \in \{(1, 2), (2, 4)\}$ . Start from  $w_0 = 1.0$ ,  $b_0 = 0.0$ , learning rate  $\eta = 0.1$ .

**Predictions.** With  $(w, b) = (1, 0)$ , predictions are  $\hat{y} = (1, 2)$ . Errors  $e = \hat{y} - y = (-1, -2)$ .

**Loss.** MSE =  $\frac{1}{2}((-1)^2 + (-2)^2) = \frac{1}{2}(1 + 4) = 2.5$ .

**Gradients.** For MSE with this batch,

$$\frac{\partial L}{\partial w} = \frac{1}{2} \sum 2e_i x_i = e_1 x_1 + e_2 x_2 = (-1) \cdot 1 + (-2) \cdot 2 = -5,$$

$$\frac{\partial L}{\partial b} = \frac{1}{2} \sum 2e_i = e_1 + e_2 = -3.$$

**Update.** Gradient descent gives

$$w_1 = w_0 - \eta \frac{\partial L}{\partial w} = 1 - 0.1 \cdot (-5) = 1.5, \quad b_1 = b_0 - \eta \frac{\partial L}{\partial b} = 0 - 0.1 \cdot (-3) = 0.3.$$

Repeating on the next mini-batch (or another epoch) will continue to reduce loss; validation loss is checked after full-dataset passes to monitor overfitting.

**Remark 9.1**

When validation loss rises while training loss falls, you may be overfitting.

## 9.2 Exercises

Why is a separate test set needed if validation loss looks good?

**Hint:** To estimate performance on truly unseen data and avoid tuning to the validation set.

---

# Chapter 10

## A Simple Neural Network Architecture

A minimal feed-forward network: inputs, a hidden layer, and an output.

### Learning Objectives

1. Sketch a tiny network with one hidden layer.
2. Track shapes through layers.
3. Explain forward pass at a high level.

### 10.1 Shapes and Flow

With  $d$  inputs,  $h$  hidden units, and one output:  $W_1 \in \mathbb{R}^{h \times d}$ ,  $\mathbf{b}_1 \in \mathbb{R}^h$ , activation  $g$ ; then  $W_2 \in \mathbb{R}^{1 \times h}$ ,  $b_2 \in \mathbb{R}$ .

**Example 10.1**

Let  $d = 2, h = 3$ . Then  $W_1$  has 3 rows and 2 columns;  $W_2$  has 1 row and 3 columns. Compute  $\mathbf{h} = g(W_1 \mathbf{x} + \mathbf{b}_1)$ , then  $\hat{y} = W_2 \mathbf{h} + b_2$ .

Take a numerical instance with ReLU activation  $g = \text{ReLU}$ :

$$\mathbf{x} = \begin{bmatrix} 2 \\ 1 \end{bmatrix}, \quad W_1 = \begin{bmatrix} 1 & -1 \\ 0.5 & 1 \\ 2 & 0 \end{bmatrix}, \quad \mathbf{b}_1 = \begin{bmatrix} 0 \\ -1 \\ 0.5 \end{bmatrix}.$$

First layer pre-activation and activation:

$$\mathbf{z}_1 = W_1 \mathbf{x} + \mathbf{b}_1 = \begin{bmatrix} 1 & -1 \\ 0.5 & 1 \\ 2 & 0 \end{bmatrix} \begin{bmatrix} 2 \\ 1 \end{bmatrix} + \begin{bmatrix} 0 \\ -1 \\ 0.5 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 4.5 \end{bmatrix}, \quad \mathbf{h} = g(\mathbf{z}_1) = \text{ReLU}(\mathbf{z}_1) = \begin{bmatrix} 1 \\ 1 \\ 4.5 \end{bmatrix}.$$

Output layer parameters and prediction:

$$W_2 = \begin{bmatrix} 1 & -2 & 0.5 \end{bmatrix}, \quad b_2 = 0.2, \quad \hat{y} = W_2 \mathbf{h} + b_2 = 1 \cdot 1 + (-2) \cdot 1 + 0.5 \cdot 4.5 + 0.2 = 1.45.$$

Thus with this small network and input  $\mathbf{x} = (2, 1)^\top$ , the hidden representation is  $\mathbf{h} = (1, 1, 4.5)^\top$  and the scalar output is  $\hat{y} = 1.45$ .

## 10.2 Exercises

If  $g = \text{ReLU}$  and  $\mathbf{h} = (0, 2, 0)^\top$ , what is  $\hat{y}$  when  $W_2 = (1, 1, 1)$  and  $b_2 = 0$ ?

**Hint:2.**

---

# Chapter 11

## Limits and Ethics

Models have limits. Responsible practice matters.

### Learning Objectives

1. Identify common failure modes: overfitting, data shift, bias.
2. Describe simple mitigations students can apply.
3. Reflect on responsible communication of model results.

### 11.1 Limits

Overfitting memorises training noise; distribution shift breaks assumptions; biased data leads to unfair outcomes.

#### Example 11.1: Overfitting

Fit a 9th-degree polynomial to 10 noisy points sampled from the straight line  $y = 2x$ . On the training points, the curve can wiggle to achieve near-zero error; on fresh test points from the same line, error is high because the wiggles chase noise rather than signal.

#### Example 11.2: Underfitting

Fit a straight line to data generated by a clear U-shaped quadratic like  $y = x^2$  with small noise. Both training and test errors remain large because a line cannot capture the curved relationship: the model is too simple for the pattern present.

### 11.2 Ethics

Use diverse, representative data when possible; report uncertainty; avoid over-claiming; consider impacts on people.

Give one reason a model trained on last year's data may underperform this year.

**Hint:** Data distribution can shift: the relationship between inputs and outputs changes.

---

# Chapter 12

## Glossary

Plain-language definitions for quick recall.

**Activation** A nonlinear function applied to a neuron's input (e.g., ReLU, sigmoid).

**Error** The difference between a prediction and the true value.

**Gradient Descent** An iterative method that adjusts parameters to reduce loss.

**Loss** A number that measures how far predictions are from targets.

**Matrix** A rectangular array of numbers; represents a linear transformation.

**Neural Network** A function built by composing linear steps with activations.

**Overfitting** Learning noise or specifics of training data that do not generalise.

**Vector** An ordered list of numbers representing features.

---

# Appendix A

## Worksheets: Checkpoints

Short comprehension checks per chapter. Replace this text with specific checkpoint prompts when ready.

---

# Appendix B

## Worksheets: Exercises

Small numeric and drawing tasks designed for practice without calculus.

---

# References

---

# Bibliography

---

# About the Author

**Vu Hung Nguyen** is the author of this book. For more information and updates:

- Website: <https://vuhung16au.github.io/>
- GitHub: <https://github.com/vuhung16au/>
- LinkedIn: <https://www.linkedin.com/in/nguyenvuhung/>